

# AWS S3

## Section introduction

- Amazon S3 is one of the main building blocks of AWS
- It's advertised as "infinitely scaling" storage
- Many websites use Amazon S3 as a backbone
- Many AWS services use Amazon S3 as an integration as well
- We'll have a step-by-step approach to S3

## Amazon S3 Use cases

- Backup and storage
- Disaster Recovery
- Archive
- Hybrid Cloud storage
- Application hosting
- Media hosting
- Data lakes & big data analytics
- Software delivery
- Static website

## Real-world examples

- **Nasdaq** stores 7 years of data into **S3 Glacier** for long-term archival
- **Sysco** runs analytics on its data stored in **S3** to gain business insights

## Amazon S3 - Buckets

- Amazon S3 allows people to store objects (files) in "buckets" (directories)
- Buckets must have a **globally unique name** (across all regions and all accounts)
- Buckets are defined at the **region level**
- S3 looks like a global service but buckets are **created in a specific region**

## Naming convention

- No uppercase letters, no underscores
- Bucket names must be 3–63 characters long
- Bucket name must not be formatted as an IP address
- Must start with a lowercase letter or number
- Must **not** start with the prefix `xn-`
- Must **not** end with the suffix `-s3alias`

## Amazon S3 - Objects

- Objects (files) have a **Key**
- The **key** is the **full path**:
  - `s3://my-bucket/my_file.txt`

- `s3://my-bucket/my_folder/another_folder/my_file.txt`
- The key is composed of **prefix** + **object name**
  - Example: `s3://my-bucket/my_folder/another_folder/my_file.txt`
- There's no concept of "directories" within buckets (although the UI might make you think otherwise)
- It's just keys with very long names that contain slashes ( / )

## Amazon S3 – Objects (cont.)

- Object values are the content of the body:
  - Max. object size is **5 TB (5000 GB)**
  - If uploading more than **5 GB**, must use **multi-part upload**
- **Metadata**: list of text key/value pairs (system or user metadata)
- **Tags**: Unicode key/value pairs (up to 10) — useful for **security** and **lifecycle policies**
- **Version ID**: present if **versioning** is enabled

## Amazon S3 – Security

### User-Based

- **IAM Policies** – define which API calls are allowed for a specific IAM user

### Resource-Based

- **Bucket Policies** – bucket-wide rules from the S3 console (allows cross-account access)
- **Object Access Control List (ACL)** – fine-grained control (can be disabled)
- **Bucket Access Control List (ACL)** – less common (can be disabled)

### Note

An **IAM principal** can access an S3 object **if**:

- The IAM permissions allow it **OR** the resource policy allows it
- **AND** there is no explicit **DENY**

### Encryption

- Encrypt objects in Amazon S3 using encryption keys

## S3 Bucket Policies

### JSON-based policies

- **Resources**: buckets and objects
- **Effect**: Allow / Deny
- **Actions**: Set of API operations to Allow or Deny
- **Principal**: The account or user to apply the policy to

### Use S3 bucket policy to:

- Grant public access to the bucket

- Force objects to be encrypted at upload
- Grant access to another account (Cross Account)

## Example: Public Access – Use Bucket Policy

An **anonymous website visitor** can access content in an **S3 bucket** if the **S3 Bucket Policy** allows **public access**.

### Key Point

- Public access is enabled via a bucket policy that explicitly allows access to everyone ( `Principal: "*"`  )
- No authentication is required for access

### Use case

- Hosting a static website with public assets (e.g., HTML, images, CSS)

### Example policy snippet

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadGetObject",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::my-public-bucket/*"
    }
  ]
}
```

## Example: User Access to S3 – IAM permissions

An **IAM User** can access an **S3 bucket** if they have an **IAM Policy** that explicitly allows access.

### Key Point

- The IAM Policy attached to the user defines which S3 actions are allowed or denied
- Access is **user-based**, controlled via IAM, not via the bucket policy

### Example use case

- Granting internal developers or employees fine-grained access to specific S3 buckets

### Example policy snippet

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
```

```

        "s3:ListBucket",
        "s3:GetObject"
    ],
    "Resource": [
        "arn:aws:s3:::example-bucket",
        "arn:aws:s3:::example-bucket/*"
    ]
}
]
}

```

## Example: EC2 instance access – Use IAM Roles

An **EC2 instance** can access an **S3 bucket**

if it has an **EC2 Instance Role** attached with the correct **IAM permissions**.

### Key Point

- The EC2 Instance Role acts as an identity for the instance
- The permissions defined in the role's policy determine what the instance can access
- No need to manage access keys manually inside the EC2 instance

### Use case

- Allowing applications running on EC2 to access S3 securely and with least privilege

### Example policy snippet (attached to the role)

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject"
      ],
      "Resource": "arn:aws:s3:::example-bucket/*"
    }
  ]
}

```

## Advanced: Cross-Account Access – Use Bucket Policy

An **IAM User from another AWS account** can access an **S3 Bucket**

if the **Bucket Policy allows cross-account access**.

### Key Point

- Cross-account access is granted by specifying the external account or user in the `Principal` field of the bucket policy
- This is a **resource-based policy** applied directly to the bucket

### Example use case

- Sharing logs, data, or backups between accounts in a multi-account AWS setup

### Example bucket policy snippet

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::123456789012:user/ExternalUser"
      },
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::example-bucket/*"
    }
  ]
}
```

## Bucket settings for Block Public Access

### Options available:

- **Block all public access:** On
- **Block public access to buckets and objects granted through new access control lists (ACLs):** On
- **Block public access to buckets and objects granted through any access control lists (ACLs):** On
- **Block public access to buckets and objects granted through new public bucket or access point policies:** On
- **Block public and cross-account access to buckets and objects through any public bucket or access point policies:** On

### Notes

- These settings were created to **prevent company data leaks**
- If you know your bucket should **never be public**, keep these **enabled**
- Settings can be applied at the **account level**

## Amazon S3 – Static Website Hosting

- S3 can host static websites and make them accessible on the Internet

### The website URL will be (depending on the region):

- `http://bucket-name.s3-website-aws-region.amazonaws.com`
- OR
- `http://bucket-name.s3-website.aws-region.amazonaws.com`

### Important

- If you get a **403 Forbidden** error, make sure the **bucket policy allows public reads**

### Example

For a bucket named `demo-bucket` in region `us-west-2` :

- `http://demo-bucket.s3-website-us-west-2.amazonaws.com`

## Amazon S3 – Versioning

- You can version your files in Amazon S3
- Versioning is **enabled at the bucket level**
- Overwriting an object with the same key generates a new **version** (e.g., 1, 2, 3...)

### Best Practices

- It's recommended to enable versioning:
  - Protects against unintended deletes (you can restore a previous version)
  - Enables easy rollback to a previous version

### Notes

- Any file that is not versioned before versioning is enabled will have version `"null"`
- Suspending versioning **does not delete previous versions**

## Amazon S3 – Replication (CRR & SRR)

- Must **enable Versioning** in both source and destination buckets
- **Cross-Region Replication (CRR)**
- **Same-Region Replication (SRR)**
- Buckets can belong to **different AWS accounts**
- Copying is **asynchronous**
- Must grant proper **IAM permissions** to S3 to perform replication

### Use Cases

- **CRR (Cross-Region Replication):**
  - Compliance
  - Lower latency access from other regions
  - Replication across AWS accounts
- **SRR (Same-Region Replication):**
  - Log aggregation
  - Live replication between production and test accounts

## Amazon S3 – Replication (Notes)

- After you enable replication, **only new objects are replicated**
- Optionally, you can replicate existing objects using **S3 Batch Replication**
  - Replicates existing objects and objects that failed replication

### For DELETE operations

- Can **replicate delete markers** from source to target (optional setting)
- Deletions with a **version ID** are *not* replicated (to avoid malicious deletes)

### No “chaining” of replication

- If bucket 1 replicates to bucket 2, and bucket 2 replicates to bucket 3,  
→ Objects created in bucket 1 are **not** replicated to bucket 3

## S3 Storage Classes

- **Amazon S3 Standard** – General purpose storage
- **Amazon S3 Standard-Infrequent Access (IA)** – For infrequently accessed data
- **Amazon S3 One Zone-Infrequent Access** – Lower-cost option for infrequent access in a single AZ
- **Amazon S3 Glacier Instant Retrieval** – Low-cost archive with milliseconds access
- **Amazon S3 Glacier Flexible Retrieval** – Archive with flexible access times (minutes to hours)
- **Amazon S3 Glacier Deep Archive** – Lowest-cost archive, access time in hours
- **Amazon S3 Intelligent Tiering** – Automatically moves data between access tiers based on usage

### Note

- Objects can move between classes **manually** or using **S3 Lifecycle configurations**

## S3 Durability and Availability

### Durability

- High durability: **99.999999999% (11 nines)** across multiple Availability Zones (AZ)
- Example: if you store **10,000,000 objects**, you can expect to lose **one object every 10,000 years**
- Durability is the **same for all S3 storage classes**

### Availability

- Measures how readily accessible the service is
- **Varies by storage class**
- Example: **S3 Standard** offers **99.99% availability**  
→ This means it may be unavailable for up to **53 minutes per year**

## S3 Standard – General Purpose

- **99.99% Availability**
- Used for **frequently accessed data**
- Provides **low latency** and **high throughput**
- Can sustain **2 concurrent facility failures**

### Use Cases

- Big Data analytics
- Mobile & gaming applications
- Content distribution

## S3 Storage Classes – Infrequent Access

- Designed for data accessed less frequently, but still requiring rapid access
- Lower cost compared to S3 Standard

### Amazon S3 Standard-Infrequent Access (S3 Standard-IA)

- **Availability:** 99.9%
- **Use cases:** Disaster Recovery, backups

### Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA)

- **Durability:** 99.999999999% (in a single AZ)
- **Availability:** 99.5%
- **Note:** Data is lost if the Availability Zone is destroyed

- **Use cases:** Secondary backup copies of on-premises data, or data that can be recreated

## Amazon S3 Glacier Storage Classes

- Low-cost object storage meant for archiving / backup
- Pricing: includes cost for storage + object retrieval

### Amazon S3 Glacier Instant Retrieval

- Millisecond retrieval, ideal for data accessed once a quarter
- Minimum storage duration: **90 days**

### Amazon S3 Glacier Flexible Retrieval (formerly Amazon S3 Glacier)

- Retrieval options:
  - Expedited: **1 to 5 minutes**
  - Standard: **3 to 5 hours**
  - Bulk: **5 to 12 hours** (free)
- Minimum storage duration: **90 days**

### Amazon S3 Glacier Deep Archive – for long-term storage

- Retrieval options:
  - Standard: **12 hours**
  - Bulk: **48 hours**
- Minimum storage duration: **180 days**

## S3 Intelligent-Tiering

- Small monthly monitoring and auto-tiering fee
- Moves objects automatically between Access Tiers based on usage
- No retrieval charges in S3 Intelligent-Tiering

### Access Tiers:

- **Frequent Access tier (automatic):** default tier
- **Infrequent Access tier (automatic):** objects not accessed for 30 days
- **Archive Instant Access tier (automatic):** objects not accessed for 90 days
- **Archive Access tier (optional):** configurable from 90 to 700+ days
- **Deep Archive Access tier (optional):** configurable from 180 to 700+ days

## S3 Storage Classes Comparison

Feature	Standard	Intelligent-Tiering	Standard-IA	One Zone-IA	Glacier Instant Retrieval	Glacier Flexible Retrieval	Glacier Deep Archive
<b>Durability</b>	99.999999999% (11 9's)						
<b>Availability</b>	99.99%	99.9%	99.9%	99.5%	99.9%	99.99%	99.99%
<b>Availability SLA</b>	99.9%	99%	99%	99%	99%	99.9%	99.9%



<b>Availability Zones</b>	>= 3	>= 3	>= 3	1	>= 3	>= 3	>= 3
<b>Min. Storage Duration</b>	None	None	30 Days	30 Days	90 Days	90 Days	180 Days
<b>Min. Billable Object Size</b>	None	None	128 KB	128 KB	128 KB	40 KB	40 KB
<b>Retrieval Fee</b>	None	None	Per GB	Per GB	Per GB	Per GB	Per GB

 [AWS Storage Classes Documentation](#)

## S3 Storage Classes – Price Comparison (us-east-1)

Feature	Standard	Intelligent-Tiering	Standard-IA	One Zone-IA	Glacier Instant Retrieval	Glacier Flexible Retrieval	Glacier Deep Archive
<b>Storage Cost (per GB)</b>	\$0.023	\$0.0025 – \$0.023	\$0.0125	\$0.01	\$0.004	\$0.0036	\$0.0009
<b>Retrieval Cost (per 1000 requests)</b>	<b>GET:</b> \$0.0004 <b>POST:</b> \$0.005	<b>GET:</b> \$0.0004 <b>POST:</b> \$0.005	<b>GET:</b> \$0.001 <b>POST:</b> \$0.01	<b>GET:</b> \$0.001 <b>POST:</b> \$0.01	<b>GET:</b> \$0.01 <b>POST:</b> \$0.02	<b>GET:</b> \$0.0004 <b>POST:</b> \$0.03 <b>Expedited:</b> \$10 <b>Standard:</b> \$0.05 <b>Bulk:</b> free	<b>GET:</b> \$0.0009 <b>POST:</b> \$0.05 <b>Standard:</b> \$0.10 <b>Bulk:</b> \$0.02
<b>Retrieval Time</b>	Instantaneous	Instantaneous				Expedited (1–5 mins) Standard (3–5 hours) Bulk (5–12 hours)	Standard (12h) Bulk
<b>Monitoring Cost</b>	–	\$0.0025	–	–	–	–	–

 [AWS S3 Pricing](#)