

Databases in AWS

Choosing the Right Database

- AWS offers many managed databases to choose from.
- To select the appropriate database, consider the following questions:
 - Is the workload read-heavy, write-heavy, or balanced?
 - What are the throughput needs? Will they fluctuate or need to scale during the day?
 - How much data do you need to store, and for how long?
 - Will the data volume grow? What is the average object size? How will the data be accessed?
 - What level of data durability is required? Is this the source of truth for your data?
 - What are the latency requirements? How many concurrent users will there be?
 - What is the data model? How will the data be queried? Are joins needed? Is the data structured or semi-structured?
 - Do you need a strong schema or more flexibility? Will the data be used for reporting or search?
 - Would a traditional RDBMS or a NoSQL solution be more appropriate?
 - Are there license costs? Can you switch to a Cloud Native DB like Aurora?

Database Types

- **RDBMS (SQL / OLTP):**
 - Amazon RDS, Amazon Aurora
 - Best suited for relational data and complex joins
- **NoSQL Databases:**
 - No joins, no SQL
 - Examples:
 - DynamoDB (~JSON structure)
 - ElastiCache (key/value store)
 - Neptune (graph database)
 - DocumentDB (compatible with MongoDB)
 - Keyspaces (compatible with Apache Cassandra)
- **Object Store:**
 - Amazon S3 (for storing large objects)
 - Amazon Glacier (for long-term backups and archives)
- **Data Warehouse (SQL Analytics / BI):**
 - Amazon Redshift (OLAP)
 - Amazon Athena
 - Amazon EMR
- **Search Engines:**
 - Amazon OpenSearch
 - Suitable for free text and unstructured data searches
- **Graph Databases:**

- Amazon Neptune
- Designed for displaying relationships between data
- **Ledger Databases:**
 - Amazon Quantum Ledger Database (QLDB)
- **Time Series Databases:**
 - Amazon Timestream

Note: Some of these databases are covered in more detail in the Data & Analytics section.

Amazon RDS – Summary

- Managed relational databases:
 - PostgreSQL, MySQL, Oracle, SQL Server, DB2, MariaDB
 - RDS Custom for Oracle & SQL Server (access to underlying instance)
- Configuration options:
 - Provisioned instance size
 - EBS volume type and size
 - Auto-scaling storage
- High availability and scalability:
 - Read Replicas
 - Multi-AZ deployments
- Security:
 - IAM integration
 - Security Groups
 - KMS (encryption at rest)
 - SSL (encryption in transit)
- Backup & recovery:
 - Automated backups (point-in-time restore, up to 35 days)
 - Manual DB snapshots for long-term backup
- Maintenance:
 - Managed and scheduled (with possible downtime)
- Authentication:
 - IAM authentication support
 - Integration with AWS Secrets Manager

Use Case:

- Store relational datasets (RDBMS / OLTP)
- Execute SQL queries and transactions

Amazon Aurora – Summary

- **Compatibility:**
 - API compatible with PostgreSQL and MySQL
 - Separation of storage and compute
- **Storage:**
 - Data stored in 6 replicas across 3 Availability Zones
 - Highly available, self-healing, auto-scaling
- **Compute:**
 - Cluster of DB instances across multiple AZs
 - Auto-scaling of read replicas
- **Cluster Architecture:**
 - Custom endpoints for writer and reader instances
- **Features:**
 - Same security, monitoring, and maintenance as Amazon RDS
 - Backup and restore options similar to RDS
 - Aurora Serverless: ideal for unpredictable or intermittent workloads, no capacity planning
 - Aurora Global: up to 16 read replicas per region, with storage replication < 1 second
 - Aurora Machine Learning: integrate with SageMaker and Comprehend
 - Aurora Database Cloning: create a new cluster faster than restoring a snapshot

Use Case:

- Same as Amazon RDS, but offers:
 - Less maintenance
 - More flexibility
 - Higher performance
 - Additional features

Amazon ElastiCache – Summary

- **Supported Engines:**
 - Managed Redis and Memcached
 - Similar to RDS but designed for caching
- **Performance:**
 - In-memory data store with sub-millisecond latency
- **Configuration:**
 - Choose an ElastiCache instance type (e.g., `cache.m6g.large`)
 - Redis supports clustering and sharding
 - Multi-AZ support with read replicas
- **Security:**
 - IAM integration
 - Security Groups
 - KMS encryption

- Redis Auth
- **Data Management:**
 - Backup and snapshot support
 - Point-in-time restore
- **Maintenance:**
 - Managed and scheduled (may involve downtime)
- **Integration:**
 - Requires changes in application code to make use of caching

Use Case:

- Key/value store
- Frequent reads, fewer writes
- Cache results of database queries
- Store session data for web applications
- No SQL support

Amazon DynamoDB – Summary

- **Overview:**
 - AWS proprietary, fully managed serverless NoSQL database
 - Millisecond latency
- **Capacity Modes:**
 - Provisioned capacity (with optional auto-scaling)
 - On-demand capacity
- **Performance & Availability:**
 - Highly available, multi-AZ by default
 - Decoupled read and write operations
 - Supports transactions
 - DAX (DynamoDB Accelerator) cluster for caching, microsecond read latency
- **Use as Key/Value Store:**
 - Can replace ElastiCache (e.g., for session data)
 - Supports TTL (Time To Live) for automatic record expiration
- **Security:**
 - IAM-based authentication and authorization
- **Event Integration:**
 - DynamoDB Streams for triggering AWS Lambda or integrating with Kinesis Data Streams
- **Global Tables:**
 - Active-active replication across regions
- **Backup and Restore:**

- Automated backups with Point-In-Time Restore (PITR) – up to 35 days
- On-demand backups
- Export to S3 without consuming read capacity units (RCU)
- Import from S3 without consuming write capacity units (WCU)
- **Schema Design:**
 - Excellent for applications that require rapid schema evolution

Use Case:

- Serverless application development
- Distributed serverless cache
- Suitable for storing small documents (hundreds of KB)

Amazon S3 – Summary

- **Storage Model:**
 - Key/value store for objects
 - Optimized for large objects (not ideal for many small ones)
- **Scalability & Limits:**
 - Serverless and infinitely scalable
 - Maximum object size: 5 TB
 - Supports versioning
- **Storage Classes (Tiers):**
 - S3 Standard
 - S3 Infrequent Access
 - S3 Intelligent Tiering
 - S3 Glacier (for archival)
 - Lifecycle policies can transition objects between tiers
- **Features:**
 - Versioning
 - Encryption
 - Replication
 - MFA-Delete
 - Access logs
- **Security:**
 - IAM policies
 - Bucket policies
 - ACLs (Access Control Lists)
 - Access Points
 - Object Lambda
 - CORS support
 - Object Lock / Vault Lock
- **Encryption Options:**
 - SSE-S3 (managed by S3)

- SSE-KMS (managed by KMS)
- SSE-C (customer-provided keys)
- Client-side encryption
- TLS for encryption in transit
- Default encryption setting
- **Operations & Automation:**
 - Batch operations with S3 Batch
 - File listing with S3 Inventory
 - Multi-part uploads
 - S3 Transfer Acceleration (improve upload/download speed)
 - S3 Select (query partial content)
 - Event Notifications via SNS, SQS, Lambda, EventBridge

Use Cases:

- Static file storage
- Key/value store for large files
- Static website hosting

Amazon DocumentDB – Summary

- **Purpose:**
 - AWS's implementation of MongoDB (just like Aurora is for PostgreSQL/MySQL)
 - Designed for storing, querying, and indexing JSON data
- **Architecture:**
 - Similar deployment concepts as Amazon Aurora
 - Fully managed and highly available
 - Replication across 3 Availability Zones
- **Storage:**
 - Automatically grows in 10 GB increments
 - Scales to support workloads with millions of requests per second

Use Case:

- Applications that use MongoDB-like data models
- NoSQL storage with JSON document support

Amazon Neptune – Summary

- **Overview:**
 - Fully managed graph database service
- **Example Use Case:**
 - Social network data model:
 - Users have friends
 - Posts have comments
 - Comments have likes from users

- Users share and like posts

- **Performance & Scalability:**

- Optimized for highly connected datasets
- Handles complex and difficult graph queries efficiently
- Millisecond latency even with billions of relationships

- **Availability:**

- Replication across 3 Availability Zones
- Up to 15 read replicas

Ideal Use Cases:

- Knowledge graphs (e.g., Wikipedia)
- Fraud detection
- Recommendation engines
- Social networking platforms

Amazon Neptune – Streams

- Provides a **real-time ordered sequence** of every change to your graph data
- Changes are available **immediately** after writing
- Ensures **no duplicates** and **strict order**
- Stream data is accessible via an **HTTP REST API**

Use Cases:

- Send notifications when specific changes occur in the graph
- Keep your graph data synchronized in another data store:
 - Amazon S3
 - Amazon OpenSearch
 - Amazon ElastiCache
- Replicate graph data across regions in Neptune

Integration Flow:

- Neptune Cluster writes data
- Neptune Streams expose changes via REST API
- A streams reader application fetches updates and forwards them to target systems

Amazon Keyspaces (for Apache Cassandra)

- **Overview:**

- Managed Apache Cassandra-compatible database service
- Based on open-source Apache Cassandra (NoSQL, distributed)

- **Features:**

- Serverless and scalable
- Fully managed and highly available
- Automatically scales tables up or down based on traffic

- Tables are replicated 3 times across multiple Availability Zones
- Uses Cassandra Query Language (CQL)
- Single-digit millisecond latency at any scale
- Supports thousands of requests per second
- **Capacity Modes:**
 - On-demand mode
 - Provisioned mode with auto-scaling
- **Data Protection:**
 - Encryption
 - Backups
 - Point-In-Time Recovery (PITR) – up to 35 days

Use Cases:

- Storing IoT device information
- Time-series data

Amazon QLDB – Summary

- **What is QLDB?**
 - QLDB stands for “Quantum Ledger Database”
 - A ledger is a system for recording financial transactions
- **Features:**
 - Fully managed and serverless
 - Highly available with replication across 3 Availability Zones
 - Designed to track and review the complete history of data changes
 - **Immutable:** entries cannot be deleted or modified
 - **Cryptographically verifiable** history of changes
- **Performance:**
 - 2–3x better performance than common blockchain ledger frameworks
 - Supports SQL-like language for data manipulation
- **Comparison with Amazon Managed Blockchain:**
 - QLDB is **not decentralized**
 - Aligned with **financial regulation compliance**

Use Cases:

- Applications that require an immutable, verifiable history of changes
- Financial transaction tracking

Amazon Timestream – Summary

- **Overview:**
 - Fully managed, fast, scalable, and serverless time series database
 - Automatically scales up or down based on demand

- Capable of storing and analyzing trillions of events per day
- **Performance & Cost:**
 - Thousands of times faster than traditional relational databases
 - Operates at 1/10th the cost
- **Features:**
 - Scheduled queries
 - Multi-measure records
 - SQL-compatible querying
 - Built-in time series analytics functions for near real-time pattern detection
- **Storage Architecture:**
 - Recent data is kept in memory
 - Historical data is moved to a cost-optimized tier
- **Security:**
 - Encryption in transit and at rest

Use Cases:

- Internet of Things (IoT) applications
- Operational monitoring and metrics
- Real-time analytics

Amazon Timestream – Architecture

Amazon Timestream integrates seamlessly with various AWS services and external tools to collect, process, analyze, and visualize time series data.

Data Ingestion Sources:

- **AWS IoT:** Collects data from IoT devices.
- **Kinesis Data Streams:** Streams large volumes of data in real time.
- **Amazon MSK (Managed Streaming for Apache Kafka):** For event-driven architectures.
- **Lambda:** Used for serverless processing and transformation of time series data.

Data Processing and Analytics:

- **Kinesis Data Analytics for Apache Flink:** Performs advanced analytics and stream processing.
- **Prometheus:** Can be integrated for collecting monitoring metrics.

Data Access and Visualization:

- **Amazon QuickSight:** For dashboarding and visualization of time series trends.
- **Amazon SageMaker:** For applying machine learning on historical time series data.
- **JDBC Connections:** External tools can connect via standard SQL-based interfaces.

Timestream acts as the central time series data store in this architecture.